



ISE-535 MODULE 2 HOMEWORK

SPRING 2025



- › For this problem, you will be using the dataset in the file named `customer_dataset.csv`
- › For each problem, you are to create the necessary Python code to answer the questions and then copy and paste screenshots of your code and the relevant response into this PowerPoint.
- › When you are ready to submit your solutions, you should save this file as a PDF file and upload it to Gradescope.
- › When uploading it to Gradescope, please follow the Gradescope prompt to indicate which pages in your solution correspond to the assignment questions



A company wants to better understand its customers by analyzing demographics, past purchase behavior, and engagement patterns. This dataset provides information that can help in customer segmentation, spending behavior analysis, and product recommendations.

Dataset variables:

Customer_ID – A unique identifier for each customer.

Age – The customer's age, useful for demographic analysis.

Number_of_Previous_Jobs – The number of jobs the customer has held, potentially indicating career stability.

Purchase_Count – Total number of purchases made by the customer.

Income_Level – Customer's income group: Low, Medium, or High.

Subscription_Years – How many years the customer has been subscribed to the company's services.

Has_Mobile_App – Whether the customer uses the mobile app (0 = No, 1 = Yes).

Annual_Spend – Total amount spent annually by the customer.

Education_Years – The number of years of education completed.

Review_Rating – The customer's average product review rating (1 to 5 stars).



- 1) Variable classification
 - › Load the data into a Pandas DataFrame
 - › Display the first few rows
 - › Classify each variable as:
 - » Measure (Continuous or Discrete)
 - » Category (Nominal or Ordinal)
 - › Modify your Pandas DataFrame as appropriate to ensure that Python treats each variable according to your classification



Screenshot of code to load dataframe and display first few rows:

Import data

```
7]: data = pd.read_csv('customer Dataset.csv')
```

Head of First Few Entries

```
9]: data.head(4)
```

| Previous_Jobs | Income_Level | Subscription_Years | Has_Mobile_App | Purchase_Count | Review_Rating | Annual_Spend | Education_Years |
|---------------|--------------|--------------------|----------------|----------------|---------------|--------------|-----------------|
| 15.000000 | High | 0.84 | 0 | 8.176449 | 5 | 44208.19 | 14.0 |
| 15.000000 | Low | 2.90 | 0 | 3.168492 | 3 | 14373.85 | NaN |
| 14.272860 | Medium | 3.57 | 1 | 6.304580 | 4 | 31689.01 | 20.0 |
| 12.412951 | Low | 1.50 | 1 | 2.328660 | 5 | 26693.14 | NaN |

Head of First Few Entries

```
] : data.head(4)
```

| | Customer_ID | Age | Number_of_Previous_Jobs | Income_Level | Subscription_Years | Has_Mobile_App | Purchase_Count | Review |
|---|-------------|-----------|-------------------------|--------------|--------------------|----------------|----------------|--------|
| 0 | 1 | 86.674826 | 15.000000 | High | 0.84 | 0 | 8.176449 | |
| 1 | 2 | 86.045678 | 15.000000 | Low | 2.90 | 0 | 3.168492 | |
| 2 | 3 | 91.364302 | 14.272860 | Medium | 3.57 | 1 | 6.304580 | |
| 3 | 4 | 92.064757 | 12.412951 | Low | 1.50 | 1 | 2.328660 | |



Variable classification table:

| Variable | Measure/Category | Continuous/Discrete Nominal/Ordinal |
|-------------------------|------------------|--|
| Customer_ID | Category | Nominal |
| Age | Measure | Discrete |
| Number_of_Previous_Jobs | Measure | Discrete |
| Purchase_Count | Measure | Discrete |
| Income_Level | Category | Ordinal |
| Subscription_Years | Measure | Discrete |
| Has_Mobile_App | Category | Nominal |
| Annual_Spend | Measure | Continuous |
| Education_Years | Measure | Discrete |
| Review_Rating | Category | Ordinal |



Make necessary modifications to your dataframe to correspond to your variable classifications. For example, if you are going to treat a numeric variable as a category, use the Pandas “astype(‘category’)” method. Display your code and the resulting datatypes (using .dtypes) on the following page

Dataset Features Conversion

```
]: ## Imputing the Null Entries
df=data
df['Age'].fillna(df['Age'].mean(), inplace=True) # type : ignore
df['Annual_Spend'].fillna(df['Annual_Spend'].median(), inplace=True) # type : ignore
df['Education_Years'].fillna(df['Education_Years'].median(), inplace=True) # type : ignore

## Converting the Data Types
df['Number_of_Previous_Jobs'] = df['Number_of_Previous_Jobs'].astype('int64') # Discrete # type : ignore
df['Has_Mobile_App'] = df['Has_Mobile_App'].astype('category') # Categorical/Nominal
df['Purchase_Count'] = df['Purchase_Count'].astype('int64') # Discrete
df['Education_Years'] = df['Education_Years'].astype('int64') # Discrete
data['Income_Level'] = data['Income_Level'].astype('category').cat.codes

## Converting Age and Subscription years

df['Age'] = df['Age'].round().astype(int)
df['Subscription_Months'] = (df['Subscription_Years'] * 12).round().astype(int)
df = df.drop(['Subscription_Years'], axis=1)

df['Income_Level'] = df['Income_Level'].astype('category')
df['Review_Rating'] = df['Review_Rating'].astype('category')
```



Updated Data Type Description

- > Dropped Identifier Column Customer ID
- > Created a new feature called subscription months and dropped subscription years for ease of interpretation

```
Age                int64
Number_of_Previous_Jobs  int64
Income_Level       category
Has_Mobile_App     category
Purchase_Count     int64
Review_Rating      category
Annual_Spend       float64
Education_Years    int64
Subscription_Months int64
dtype: object
```

```
[361]: data.head(6)
```

```
[361]:
```

| | Age | Number_of_Previous_Jobs | Income_Level | Has_Mobile_App | Purchase_Count | Review_Rating | Annual_Spend | Education_Years | Subscription_Months |
|---|-----|-------------------------|--------------|----------------|----------------|---------------|--------------|-----------------|---------------------|
| 0 | 87 | 15 | 0 | 0 | 8 | 5 | 44208.19 | 14 | 10 |
| 1 | 86 | 15 | 1 | 0 | 3 | 3 | 14373.85 | 13 | 35 |
| 2 | 91 | 14 | 2 | 1 | 6 | 4 | 31689.01 | 20 | 43 |
| 3 | 92 | 12 | 1 | 1 | 2 | 5 | 26693.14 | 13 | 18 |
| 4 | 85 | 12 | 1 | 0 | 4 | 3 | 12150.15 | 13 | 23 |
| 5 | 94 | 15 | 1 | 0 | 1 | 3 | 18146.82 | 13 | 5 |



2) Univariate Analysis – Descriptive Statistics

Create a table summarizing the descriptive statistics for your measures. On the following page, paste a screenshot of your code to generate the table and the table itself. Be sure that your statistics include Skew and Kurtosis



```
[367... import pandas as pd
import scipy.stats as stats

numeric_columns = data.select_dtypes(include=['number']).columns
desc_stats = data[numeric_columns].describe().T

# Calculate skewness and kurtosis for numeric columns
desc_stats['skew'] = data[numeric_columns].skew()
desc_stats['kurtosis'] = data[numeric_columns].kurtosis()

# For categorical columns, unique values and mode
for col in data.select_dtypes(include=['category']).columns:
    desc_stats.loc[col, 'unique'] = data[col].nunique()
    desc_stats.loc[col, 'mode'] = data[col].mode()[0]

desc_stats
```

| | count | mean | std | min | 25% | 50% | 75% | max | skew | kurtosis | unique | mode |
|-------------------------|--------|-------------|-------------|----------|-----------|----------|-----------|----------|----------|-----------|--------|------|
| Age | 1000.0 | 35.66500 | 7.214348 | 19.00 | 32.0000 | 35.000 | 38.000 | 94.00 | 4.187335 | 29.986393 | NaN | NaN |
| Number_of_Previous_Jobs | 1000.0 | 2.68400 | 2.020452 | 0.00 | 1.0000 | 2.000 | 4.000 | 15.00 | 1.931317 | 8.550344 | NaN | NaN |
| Purchase_Count | 1000.0 | 3.45000 | 2.033134 | 0.00 | 2.0000 | 3.000 | 5.000 | 12.00 | 0.499586 | 0.298372 | NaN | NaN |
| Annual_Spend | 1000.0 | 11433.16554 | 8542.161652 | -4798.32 | 5506.7375 | 9464.785 | 15562.355 | 44208.19 | 1.073247 | 0.786004 | NaN | NaN |
| Education_Years | 1000.0 | 13.04100 | 2.886807 | 6.00 | 13.0000 | 13.000 | 13.000 | 20.00 | 0.037561 | 1.043922 | NaN | NaN |
| Subscription_Months | 1000.0 | 59.91000 | 60.908339 | 0.00 | 17.0000 | 41.000 | 81.250 | 452.00 | 1.991001 | 5.528214 | NaN | NaN |
| Income_Level | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 3.0 | 1.0 |
| Has_Mobile_App | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 2.0 | 1.0 |
| Review_Rating | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 5.0 | 5.0 |



2) Univariate Analysis – Measure Visualizations

Create a series of histograms and boxplots to visualize your measures. On the following pages, paste screenshots of your code and of the visualizations.

Univariate measure Visualizations

```
371]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(15, 10))

# List of numeric columns for visualization
numeric_columns = data.select_dtypes(include=['number']).columns

# Generate histograms
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(2, len(numeric_columns)//2, i)
    sns.histplot(data[col], kde=True)
    plt.title(f'Histogram of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')

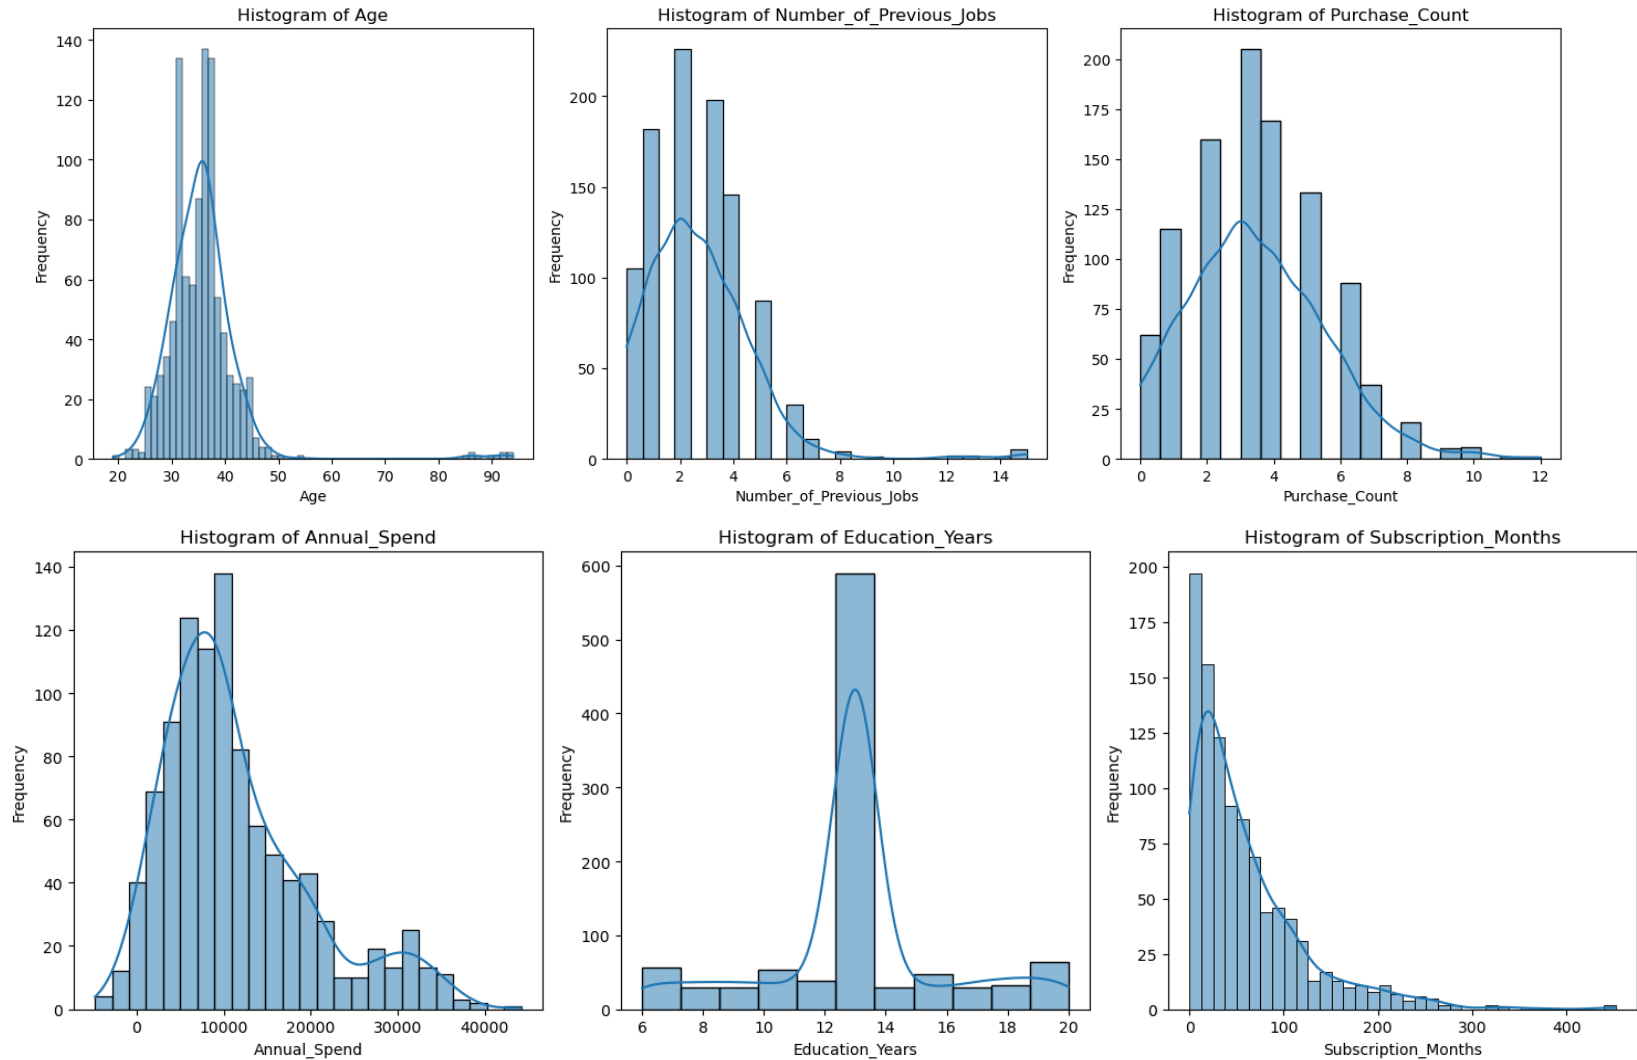
plt.tight_layout()
plt.show()

# Generate boxplots
plt.figure(figsize=(15, 10))
for i, col in enumerate(numeric_columns, 1):
    plt.subplot(2, len(numeric_columns)//2, i)
    sns.boxplot(x=data[col])
    plt.title(f'Boxplot of {col}')
    plt.xlabel(col)

plt.tight_layout()
plt.show()
```

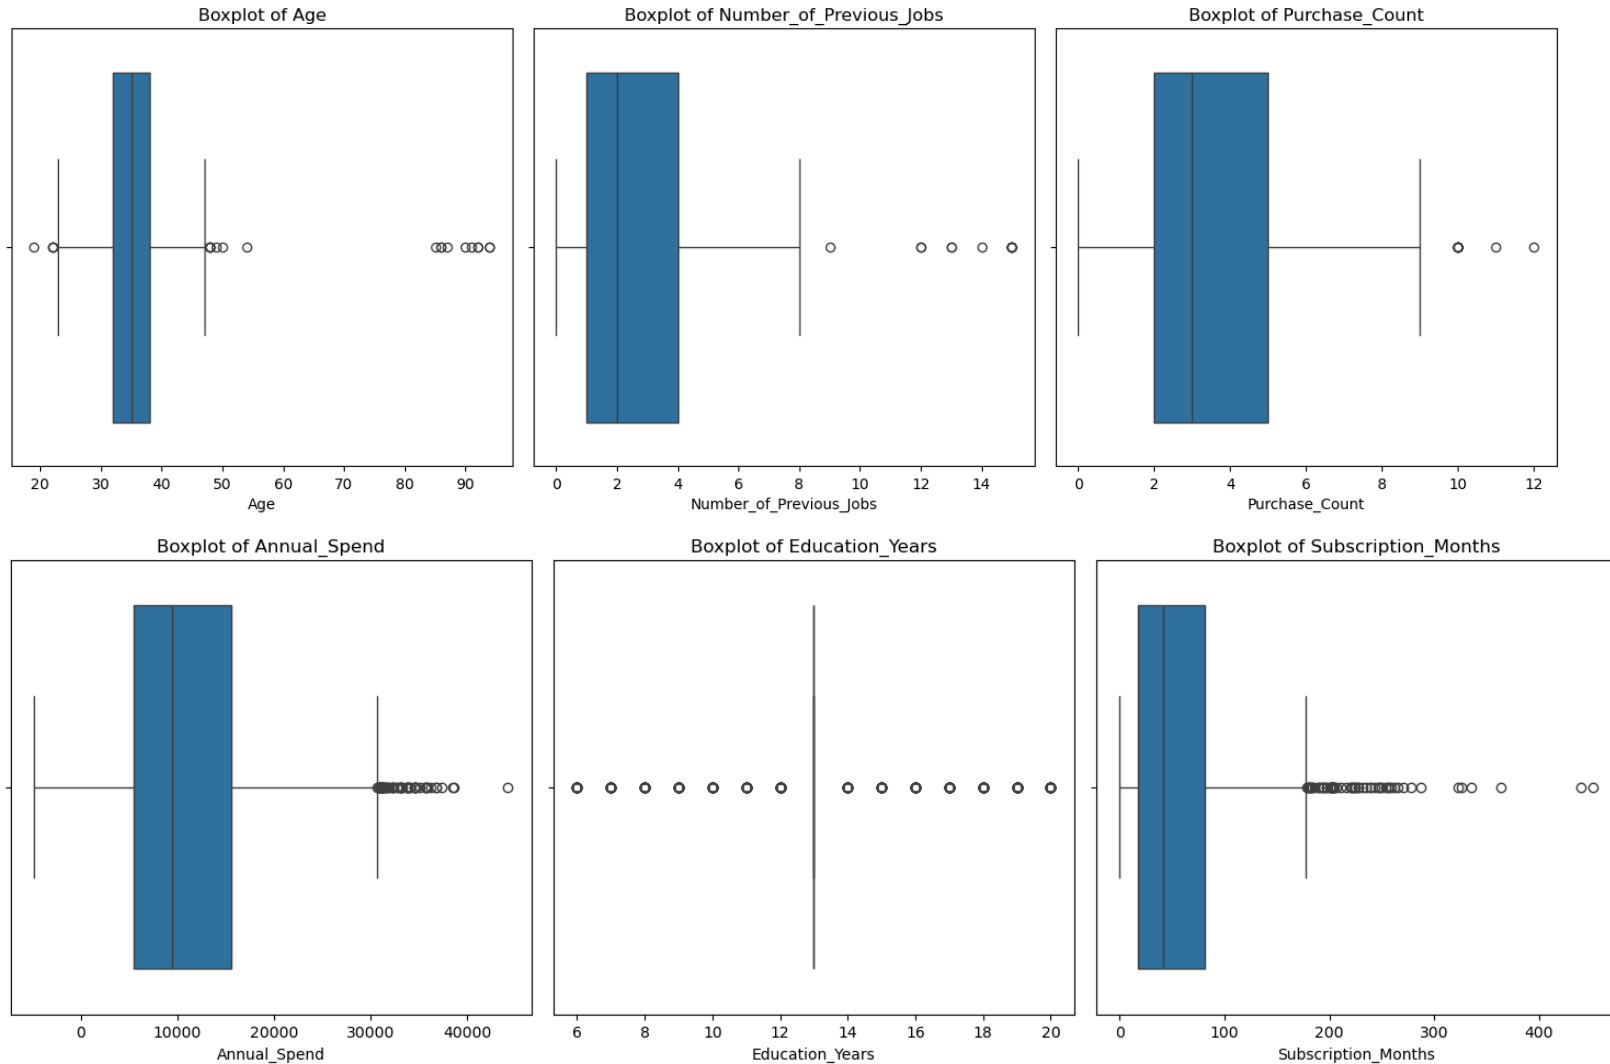


2) Univariate Analysis – Measure Visualizations





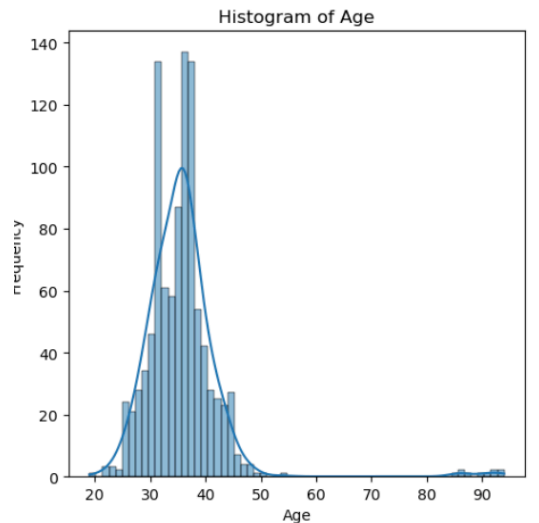
2) Univariate Analysis – Measure Visualizations





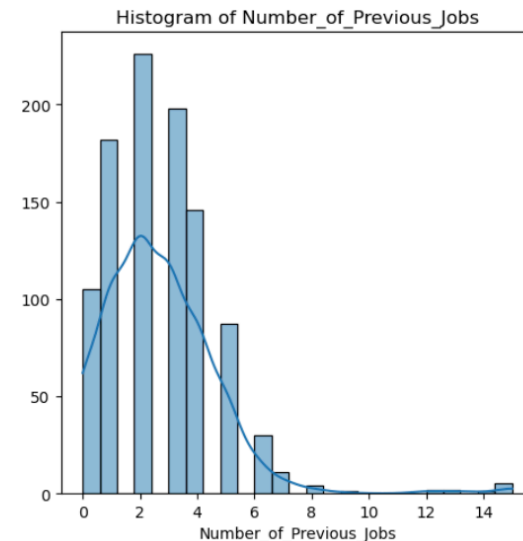
2) Univariate Analysis – Measures

Based on a review of the histograms comment below on the distribution shape of each variable and whether it seems reasonable given the nature of the variable



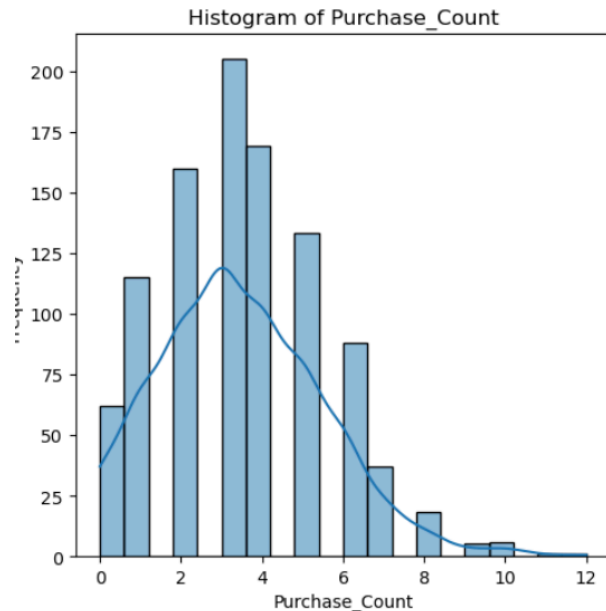
Age : Shows a distribution with highly right skewed values including outliers to the right with age group of 80-90+. The distribution is peaked up and not widely spread out, having a high kurtosis. Seems reasonable.

Number of Previous Jobs: It is right skewed with some outliers having 10-14 jobs; with moderate kurtosis with mean values of previous jobs as 2.68 (close to 3). Seems reasonable but has outliers with 10-14 previous jobs



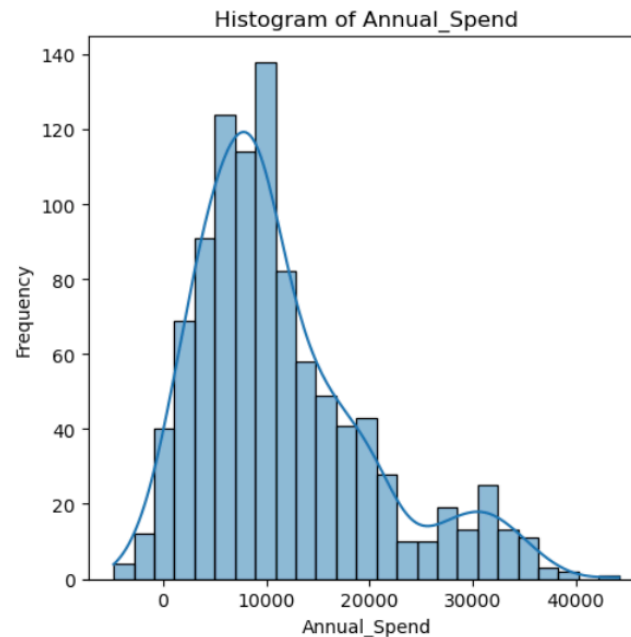


2) Univariate Analysis – Measure Visualizations



Purchase Count: The distribution is slightly right skewed with an average purchase made by customers as 3.4 (3-4). Seems Reasonable for the variable

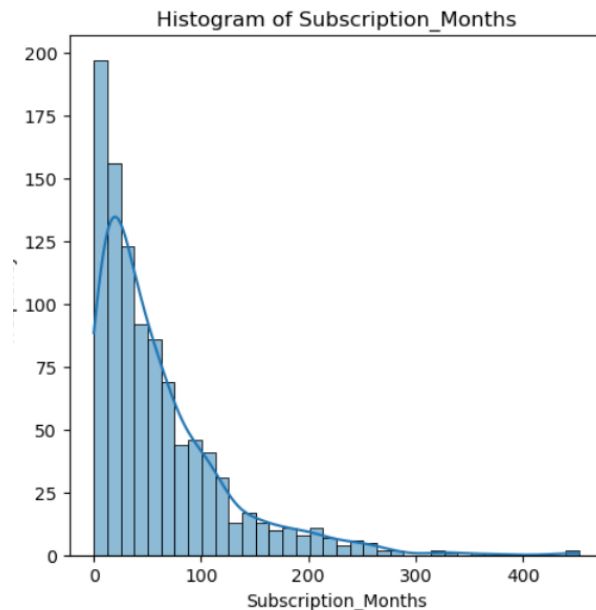
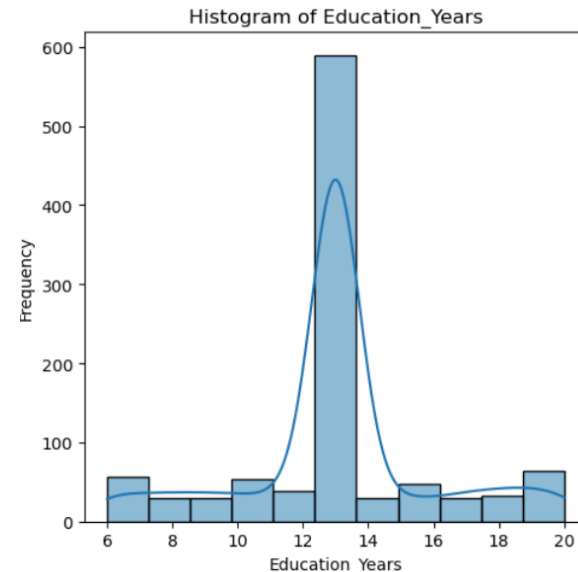
Annual Spend : The distribution is right skewed and seems reasonable with average spending amount around \$11,433, having outliers with people making high value purchases. Seems Reasonable, and depends.





2) Univariate Analysis – Measure Visualizations

Education years : Seems Reasonable with most people having around 12 years of education with very light skewness - resembling a normal distribution. Almost equal amount of people have 6-12 and 12-20 years of education.



Subscription Months: Perfectly right skewed having outliers of 200+ months of subscription and majority of customers being not subscribed. Seems Reasonable.



2) Univariate Analysis – Measures

Based on a review of the histograms and box plot, which variable appears to have the most extreme outliers?

- › Age has Most Extreme outliers with high skewness of 4.187335



2) Univariate Analysis – Categories

Create bar charts of the categories (don't include Customer_ID). Paste on the following pages your code and the bar chart visualizations

```
76]: import matplotlib.pyplot as plt
import seaborn as sns
import warnings


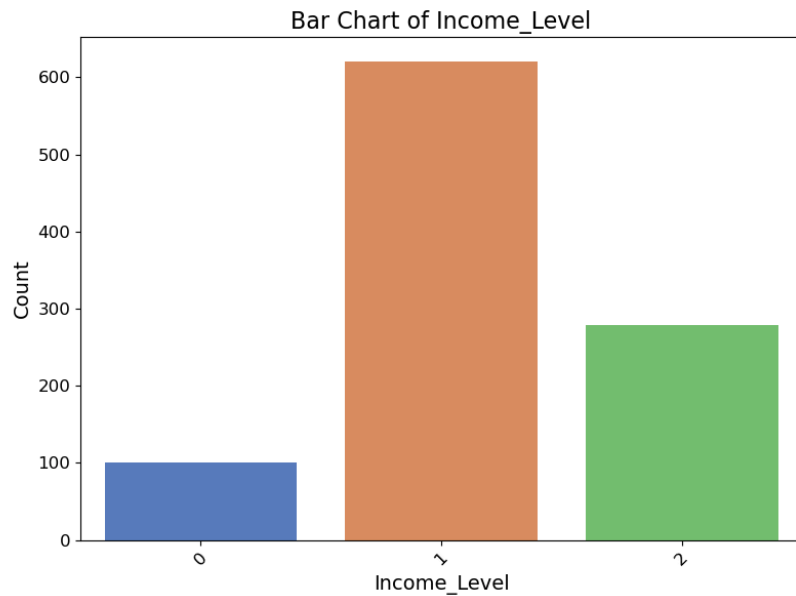

# Ignore warnings
warnings.filterwarnings('ignore')

# List of categorical columns to plot
categorical_columns = ['Income_Level', 'Has_Mobile_App', 'Review_Rating']

# Create a bar plot for each categorical column
for col in categorical_columns:
    plt.figure(figsize=(8, 6)) # Set figure size
    sns.countplot(data=data, x=col, palette='muted')
    plt.title(f'Bar Chart of {col}', fontsize=16)
    plt.xlabel(col, fontsize=14) # X-axis label
    plt.ylabel('Count', fontsize=14) # Y-axis label
    plt.xticks(rotation=45, fontsize=12)
    plt.yticks(fontsize=12)
    plt.tight_layout()
    plt.show() # Show plot
```



2) Univariate Analysis – Categories





3) Univariate Analysis – Diagnosis. Paste below a table listing the number of missing values in each variable

Missing Data Handling

```
[322]: data.isna().sum()
```

```
[322]: Customer_ID      0
      Age              60
      Number_of_Previous_Jobs  0
      Income_Level      0
      Subscription_Years  0
      Has_Mobile_App     0
      Purchase_Count     0
      Review_Rating      0
      Annual_Spend       40
      Education_Years    558
      dtype: int64
```



3) For each variable, we wish to determine if they are Missing Completely at Random (MCAR) or Missing at Random (MAR). A simple test is to create a new binary column in the dataframe for each variable that has missing values that indicates whether that row is missing that variable. Then, using this new column, calculate frequency counts for each of the other category variables for those rows with and without missing values. If the frequency counts are significantly different, that is an indication that variable may not be MCAR.

On the following pages, show your code that does this analysis and summarize your conclusions.



Missing Value Analysis

```
[392]: data_m = pd.read_csv('customer Dataset.csv')

# Creating binary columns indicating whether each variable is missing
for col in ['Age', 'Annual_Spend', 'Education_Years']:
    data_m[f'{col}_missing'] = data_m[col].isna().astype(int)

categorical_columns = ['Income_Level', 'Has_Mobile_App', 'Review_Rating']

# For each variable with missing data, calculate frequency counts by 'missing' status
for col in ['Age', 'Annual_Spend', 'Education_Years']:
    print(f"\nFrequency counts for missing vs non-missing {col} values:\n")
    for cat_col in categorical_columns:
        print(f"\n--- {cat_col} ---")
        print("Missing:")
        print(data_m[data_m[f'{col}_missing'] == 1][cat_col].value_counts())
        print("Not Missing:")
        print(data_m[data_m[f'{col}_missing'] == 0][cat_col].value_counts())
    print("\n")
```



Variables Age, Annual Spending, Education years has missing values

Analyzing the Missing Criteria in Age:

```
--- Income_Level ---
Missing:
Income_Level
Low      43
Medium   11
High      6
Name: count, dtype: int64
Not Missing:
Income_Level
Low      578
Medium   268
High      94
Name: count, dtype: int64
```

```
--- Has_Mobile_App ---
Missing:
Has_Mobile_App
1      30
0      30
Name: count, dtype: int64
Not Missing:
Has_Mobile_App
1      472
0      468
Name: count, dtype: int64
```

```
--- Review_Rating ---
Missing:
Review_Rating
5      16
4      14
3      12
1      10
2       8
Name: count, dtype: int64
Not Missing:
Review_Rating
5      250
3      238
4      197
2      156
1       99
Name: count, dtype: int64
```



Analyzing the Missing Criteria in Annual Spending:

--- Income_Level ---

Missing:

Income_Level

Low 24

Medium 13

High 3

Name: count, dtype: int64

Not Missing:

Income_Level

Low 597

Medium 266

High 97

Name: count, dtype: int64

--- Has_Mobile_App ---

Missing:

Has_Mobile_App

0 22

1 18

Name: count, dtype: int64

Not Missing:

Has_Mobile_App

1 484

0 476

Name: count, dtype: int64

--- Review_Rating ---

Missing:

Review_Rating

5 17

4 10

3 7

2 4

1 2

Name: count, dtype: int64

Not Missing:

Review_Rating

5 249

3 243

4 201

2 160

1 107

Name: count, dtype: int64



Analyzing the Missing Criteria in Education Years:

```
--- Income_Level ---
Missing:
Income_Level
Low      558
Name: count, dtype: int64
Not Missing:
Income_Level
Medium   279
High     100
Low       63
Name: count, dtype: int64
```

```
--- Has_Mobile_App ---
Missing:
Has_Mobile_App
0        279
1        279
Name: count, dtype: int64
Not Missing:
Has_Mobile_App
1        223
0        219
Name: count, dtype: int64
```

```
--- Review_Rating ---
Missing:
Review_Rating
5        157
3        146
4        119
2         84
1         52
Name: count, dtype: int64
Not Missing:
Review_Rating
5        109
3        104
4         92
2         80
1         57
Name: count, dtype: int64
```



Conclusions:

> * **Missing in Age : (MCAR)**

1. By looking at the Income level, the age variable seems to be missing for low income group, however, the not missing values shows otherwise, considering the proportion size of each category.
2. Having a mobile app does not seem to have any relation with missingness in Age.
3. Missing Ages are high in high reviews but also considering the Not missing values ranges of review ratings, it shows that the age is randomly missing without any underlying strong evidence. Hence, this could be Missing Completely at Random. (MCAR)

* **Missing in Annual Spending : (MAR)**

1. The low income level group has a high missing reports of annual spend, but the not missing stats show this might be at random.
2. Having a mobile app does seem to have a slight relation with missingness in Annual spending, with those without apps having no annual spending reported, whereas those with apps having less missing values.
3. Missing Annual Spendings are high in high reviews but also considering the Not missing values ranges of review ratings, it shows that the age is randomly missing without any underlying strong evidence. This could be Missing at Random. (MAR)



*** Missing in Education years : (MNCAR) - Not missing completely at Random**

1. By looking at the Income level, the Education years variable seems to be missing one and only for low income group. This is a strong indication that customers fail to provide education years whose income levels are low.
2. Having a mobile app does not seem to have any relation with missingness in Education Years.
3. Missing education levels are high in high reviews but also considering the Not missing values ranges of review ratings, it shows that the education is randomly missing without any underlying strong evidence with respect to this.

The variable Education can be considered not missing at random considering the Income Level.



4) Create and display below a correlation matrix of the measures in your dataset


```
[66]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Assuming 'data' is the dataframe containing your dataset

# Step 1: Select numeric variables (you can adjust this list based on your dataset)
numeric_columns = ['Age', 'Income_Level', 'Subscription_Months', 'Purchase_Count', 'Annual_Spend', 'Education_Years', 'Number_of_Previous_Jobs']

# Step 2: Calculate the correlation matrix
correlation_matrix = data[numeric_columns].corr()

# Step 3: Visualize the correlation matrix as a heatmap
plt.figure(figsize=(10, 8)) # Set the size of the figure
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', cbar=True, square=True, linewidths=0.5)
plt.title("Correlation Matrix of Numeric Variables")
plt.show()
```





4) Comment below on the correlation matrix. What correlations exist? Do you see anything unexpected?

Weak Positive Correlation: Age and Annual Spend (0.15)

This makes sense as Young people (teens to young adults) might spend slightly less than middle aged or adults who might be more financially stable.




Moderate Positive Correlation : Annual Spend and Purchase Count (0.35)

This also makes sense as more the annual expense more the number of purchases made.

High Positive Correlation: 0.71 Correlation between Age and Number of previous jobs



4) Create and display below side-by-side boxplots of each category against Annual_Spend.





4) Comment below on the side-by-side boxplots. What correlations exist? Do you see anything unexpected?

The Income Level and Annual spending, has an unexpected behavior having more annual spending for low income category customers, whereas less annual spendings for high income level compared to low income customers. We might assume there might be a mindless spending for low income customers, compared to medium-high income cohort showing a diligent spending.

The presence of mobile app and Annual spending does not have any significant correlation.

The annual spending of high review customers is higher than low reviewed customers, and shows a gradual progression of spendings with review rates, which is quite expected.



4) Create scatterplots of each measure vs. annual spend. Paste your code and resulting plots on the following page(s).

Measures Vs Annual Spend

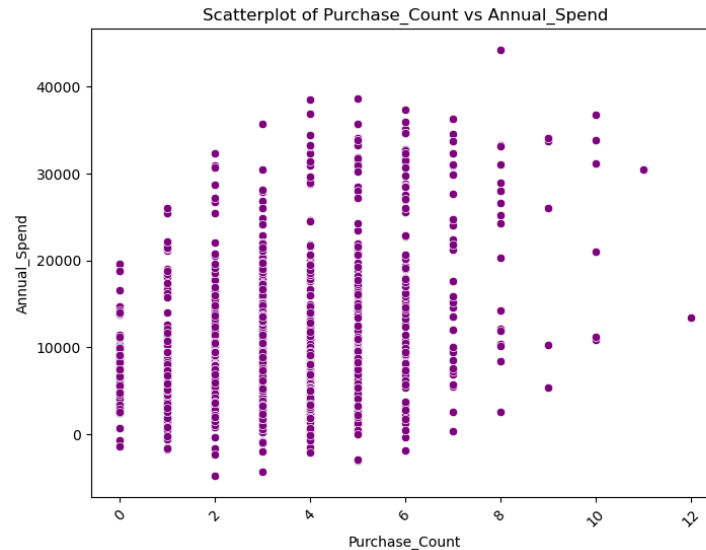
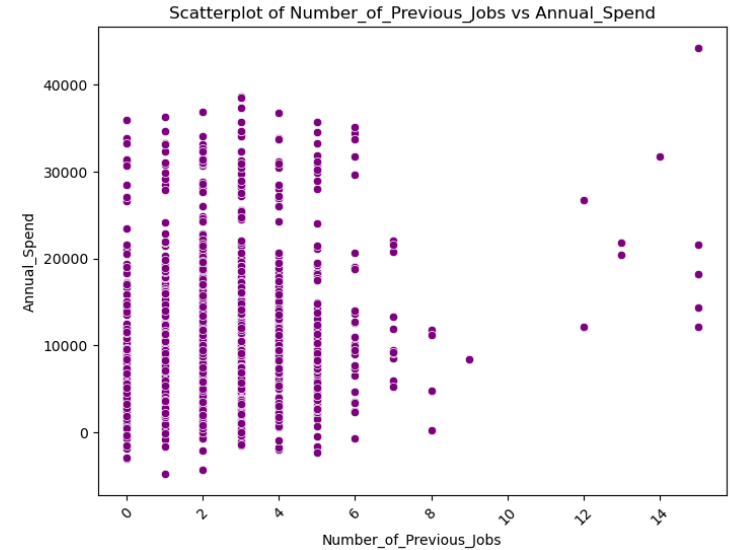
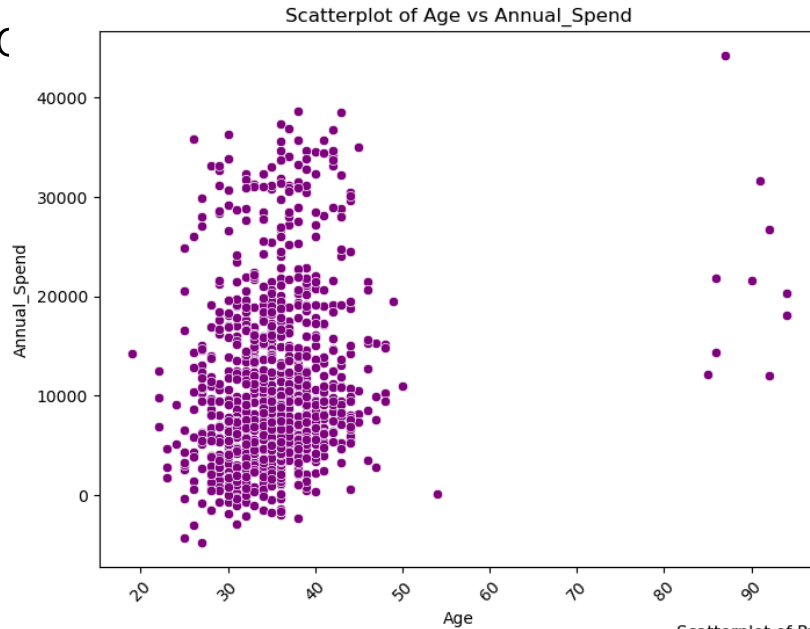
```
.8]: import seaborn as sns
import matplotlib.pyplot as plt

numerical_vars = ['Age', 'Number_of_Previous_Jobs', 'Purchase_Count', 'Education_Years', 'Subscription_Months']

# Step 2: Create scatterplot for each numerical variable against Annual_Spend
for var in numerical_vars:
    plt.figure(figsize=(8, 6)) # Set the size of the figure
    sns.scatterplot(x=var, y='Annual_Spend', data=data, color='purple')
    plt.title(f'Scatterplot of {var} vs Annual_Spend')
    plt.xticks(rotation=45) # Rotate x-axis labels if necessary
    plt.show()
```





Scatterplot






Scatterplots

Scatterplot of Education_Years vs Annual_Spend



Scatterplot of Subscription_Months vs Annual_Spend





5) To identify potential multicollinearity, calculate the Variance Inflation Factor (VIF) statistics for each of the measures. Hint: use the `variance_inflation_factor` function in the `statsmodels.stats.outlier_influence` library. On the following code display your code and the resulting statistics.

```
[70]: import pandas as pd
      from statsmodels.stats.outliers_influence import variance_inflation_factor
      from statsmodels.tools.tools import add_constant

      numerical_vars = ['Age', 'Number_of_Previous_Jobs', 'Purchase_Count',
                        'Annual_Spend', 'Education_Years', 'Subscription_Months']
      X = data[numerical_vars]

      vif_data = pd.DataFrame()
      vif_data["variable"] = X.columns
      vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

      print(vif_data)
```



VIF calculations:

| | variable | VIF |
|---|-------------------------|-----------|
| 0 | Age | 25.081613 |
| 1 | Number_of_Previous_Jobs | 4.765390 |
| 2 | Purchase_Count | 4.309726 |
| 3 | Annual_Spend | 3.314980 |
| 4 | Education_Years | 14.715171 |
| 5 | Subscription_Months | 1.943472 |



VIF calculations: Comment below on your interpretation of the VIF statistics

Age: Age has a high VIF of 25.08 showing high significance to the other variables, especially the number of years of education.

Number of education years: A VIF 14.715 shows a high significance among the other variables.

Since the age and number of years of education has a high variance compared to others, they are considered to be highly significant and potentially signify multicollinearity.